

Think Like The Cardinality Estimator

Taio b Ali



Data Solutions Manager, GMO LLC



<http://sqlworldwide.com/>



/sqlworldwide



@sqlworldwide



Taiob@sqlworldwide.com

Data Professional

Microsoft Data Platform MVP. 14 Years working with Microsoft Data Platform. Microsoft and MongoDB certified. Worked in ecommerce, healthcare and finance industry.

Giving Back

Board member NESQL user group and PASS DBA virtual group. Organizer of Boston SQL Saturday. Frequent speaker at local and virtual user groups, SQL Saturdays, and azure events.

When Not Working

Running – One 26.2 and many 13.1 miles. Learning US history. Shuttling 3 kids.

Predicate, Density, Predicate Selectivity

What is Cardinality

Why Cardinality Matters

DBCC SHOW_STATISTICS

Magic Numbers

Predicate

Is an expression that evaluates to
TRUE, FALSE, or UNKNOWN

- Join
- Filter
 - Where
 - Having

```
SELECT
    cus.CustomerID,
    COUNT(0) AS [NumOfOrders]
FROM
    sales.Orders AS ord
JOIN
    sales.Customers AS cus
ON
    ord.CustomerID=cus.CustomerID
WHERE
    ord.OrderDate='2013-01-01'
GROUP BY cus.CustomerID
HAVING COUNT(0) > 2 ;
GO
```

Join Predicate


Filter Predicate
Where/Having

Density

How often duplicate values occur in a column.

$$\frac{1}{[\text{\# of distinct values in a column}]}$$

```
SELECT  
    COUNT(DISTINCT customerID) AS [DistinctCusId]  
FROM  
    sales.Orders;  
GO
```



Density
663 distinct customerID
 $1/663=0.00150830$

Predicate Selectivity

Most commonly used to describe a predicate.

$$\frac{[\text{\# rows that pass the predicate}]}{[\text{total number of rows}]}$$

```
SELECT
    COUNT(0) AS [NumOfOrders]
FROM
    sales.Orders
WHERE
    CustomerID=577;
GO
```

75 rows for
customerID 577
73595 total number
of rows.
 $75/73595=.0010$



Cardinality

For Us

- Number of rows returned by a query operator.

Structured Query Language

- Uniqueness of data values contained in a particular column (attribute) of a database table.

Math

- Cardinality of a set is a measure of the “number of elements of the set”.

Clustered Index Seek (Clustered)

Scanning a particular range of rows from a clustered index.

Physical Operation	Clustered Index Seek
Logical Operation	Clustered Index Seek
Actual Execution Mode	Row
Estimated Execution Mode	Row
Storage	RowStore
Actual Number of Rows	12
Actual Number of Batches	0
Estimated Operator Cost	0.0032904 (8%)
Estimated I/O Cost	0.003125
Estimated Subtree Cost	0.0032904
Estimated CPU Cost	0.0001654
Estimated Number of Executions	1
Number of Executions	1
Estimated Number of Rows	7.61577
Estimated Row Size	143 B

Actual



Estimated



Why Cardinality Matters?

Before

**Parse
+
Normalize**

**Cardinality
Estimation**

**Derive
Cardinality
+
Create
Many Plans**

Downstream

Optimization

Downstream

**Execute
Plan**



Why Cardinality Matters?

Cost

Parallel

Serial

Memory Grant

In Memory

Spill to Disk

Access Method

Seek

Scan

Seek + Scan

Algorithm

Join

Aggregate

Sort

DBCC SHOW_STATISTICS

Header

- Meta data about the statistics.

Density Vector

- How many unique values are present within a column or columns?

Histogram

- Frequency of data within the first key column of the statistics.

- STAT_HEADER

Name	Updated	Rows	Rows Sampled	Steps	Density	Average key length	String Index	Filter Expression	Unfiltered Rows
FK_Sales_Orders_ContactPersonID	Jun 2 2016 10:45AM	71583	71583	200	0.009388237	8	NO	NULL	71583

- DENSITY_VECTOR

	All density	Average Length	Columns
1	0.001512859	4	ContactPersonID
2	1.39698E-05	8	ContactPersonID, OrderID

- HISTOGRAM

	RANGE_HI_KEY	RANGE_ROWS	EQ_ROWS	DISTINCT_RANGE_ROWS	AVG_RANGE_ROWS
1	1001	0	128	0	1
2	1007	248	107	2	124
3	1013	231	130	2	115.5
4	1021	327	111	3	109
5	1025	101	85	1	101
6	1031	237	133	2	118.5

STAT_HEADER

Name	Updated	Rows	Rows Sampled	Steps
NCI_FilteredContactPersonID	Mar 31 2017 4:55PM	50299	50299	200

Deprecated

Density	Average key length	String Index	Filter Expression	Unfiltered Rows
0.00945746	8	NO	[(contactpersonid)>(2000))	73595

DENSITY_VECTOR

1/Number of
distinct values in
column


Names of columns
in the prefix

All density	Average Length	Columns
0.002164502	4	ContactPersonID
1.988111E-05	8	ContactPersonID, OrderID


Histogram

	RANGE_HI_KEY	RANGE_ROWS	EQ_ROWS	DISTINCT_RANGE_ROWS	AVG_RANGE_ROWS
70	2035	330	123	3	110
71	2043	333	90	3	111
72	2051	317	94	3	105.6667
73	2055	106	114	1	106
74	2061	249	122	2	124.5
75	2065	107	108	1	107
76	2073	326	111	3	108.6667
77	2077	127	129	1	127
78	2083	234	119	2	117
79	2091	332	106	3	110.6667
80	2095	118	135	1	118
81	2103	338	125	3	112.6667
82	2107	112	110	1	112
83	2113	224	109	2	112
84	2121	375	118	3	125
85	2127	224	101	2	112

RANGE_HI_KEY	RANGE_ROWS	EQ_ROWS	DISTINCT_RANGE_ROWS	AVG_RANGE_ROWS
2083		119		
Between 2084 and 2090	332		3	110.66
2091		106		



- Single Predicate
 - Histogram direct hit
 - Histogram intra step
 - Scaling
 - Distinct



- Multiple Predicates
 - Conjunction
 - Disjunction
- Parameter Sniffing
- Unknown
- Ascending Key



SQL2019 CU5
SSMS 18.5.1

Resource

- [Query Tuning Fundamentals](#)
- [DBCC SHOW_STATISTICS \(Transact-SQL\)](#)
- [13 Things You Should Know About Statistics and the Query Optimizer](#)
- [Cardinality Estimation for Multiple Predicates](#)
- [New Trace Flag to Fix Table Variable Performance](#)
- [Ascending key Issue TF 2389 and 2390](#)
- [Optimizing Query Plans with the SQL Server 2014 Cardinality Estimator](#)
- [Cardinality Estimation \(SQL Server\)](#)



@SqlWorldWide



linkedin.com/in/sqlworldwide



sqlworldwide.com



taio@sqlworldwide.com

